

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.impute import SimpleImputer
from scipy.stats import zscore
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.decomposition import PCA
from sklearn import metrics
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
```

In [2]:

```
veh= pd.read_csv("vehicle-1.csv") # Loading the dataset.
```

In [3]:

```
veh.shape # shape od the dataset.
```

Out[3]:

```
(846, 19)
```

In [4]:

```
veh.head() # to get a general idea of data.
```

Out[4]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_aspect_ratio	scatter_ratio	elongatedness
0	95	48.0	83.0	178.0	72.0	10	162.0	42.0
1	91	41.0	84.0	141.0	57.0	9	149.0	45.0
2	104	50.0	106.0	209.0	66.0	10	207.0	32.0
3	93	41.0	82.0	159.0	63.0	9	144.0	46.0
4	85	44.0	70.0	205.0	103.0	52	149.0	45.0

In [5]:

```
veh.info() #Identifying datatype of each columns.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 846 entries, 0 to 845
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   compactness                          846 non-null    int64
1   circularity                          841 non-null    float64
2   distance_circularity                 842 non-null    float64
3   radius_ratio                        840 non-null    float64
4   pr.axis_aspect_ratio                844 non-null    float64
5   max.length_aspect_ratio             846 non-null    int64
6   scatter_ratio                       845 non-null    float64
7   elongatedness                       845 non-null    float64
8   pr.axis_rectangularity              843 non-null    float64
9   max.length_rectangularity           846 non-null    int64
10  scaled_variance                     843 non-null    float64
11  scaled_variance.1                   844 non-null    float64
12  scaled_radius_of_gyration           844 non-null    float64
```

```

13 scaled_radius_of_gyration.1  842 non-null    float64
14 skewness_about               840 non-null    float64
15 skewness_about.1             845 non-null    float64
16 skewness_about.2             845 non-null    float64
17 hollows_ratio                846 non-null    int64
18 class                        846 non-null    object

```

dtypes: float64(14), int64(4), object(1)

memory usage: 125.7+ KB

In [6]:

```
veh.describe()
```

Out[6]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_aspect_ratio	scatter_ratio	elongate
count	846.000000	841.000000	842.000000	840.000000	844.000000	846.000000	845.000000	845.000000
mean	93.678487	44.828775	82.110451	168.888095	61.678910	8.567376	168.901775	40.901775
std	8.234474	6.152172	15.778292	33.520198	7.891463	4.601217	33.214848	7.891463
min	73.000000	33.000000	40.000000	104.000000	47.000000	2.000000	112.000000	26.000000
25%	87.000000	40.000000	70.000000	141.000000	57.000000	7.000000	147.000000	33.000000
50%	93.000000	44.000000	80.000000	167.000000	61.000000	8.000000	157.000000	43.000000
75%	100.000000	49.000000	98.000000	195.000000	65.000000	10.000000	198.000000	46.000000
max	119.000000	59.000000	112.000000	333.000000	138.000000	55.000000	265.000000	61.000000

In [7]:

```
veh.isnull().values.any() # To check if there are null values.
```

Out[7]:

True

In [8]:

```

labelencoder = LabelEncoder()
veh['class_dup'] = labelencoder.fit_transform(veh['class'])
veh

```

Out[8]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_aspect_ratio	scatter_ratio	elongateddn
0	95	48.0	83.0	178.0	72.0	10	162.0	4
1	91	41.0	84.0	141.0	57.0	9	149.0	4
2	104	50.0	106.0	209.0	66.0	10	207.0	3
3	93	41.0	82.0	159.0	63.0	9	144.0	4
4	85	44.0	70.0	205.0	103.0	52	149.0	4
...
841	93	39.0	87.0	183.0	64.0	8	169.0	4
842	89	46.0	84.0	163.0	66.0	11	159.0	4
843	106	54.0	101.0	222.0	67.0	12	222.0	3
844	86	36.0	78.0	146.0	58.0	7	135.0	5
845	85	36.0	66.0	123.0	55.0	5	120.0	5

846 rows × 20 columns

In the above step, The class was a categorical column. \ This can't be used in SVM algorithms. Thus the columns elements were given number. \ Such as bus is 0, Van is 2 and car is 1.

In [9]:

```
veh.mean()
```

Out[9]:

```
compactness           93.678487
circularity           44.828775
distance_circularity  82.110451
radius_ratio         168.888095
pr.axis_aspect_ratio  61.678910
max.length_aspect_ratio  8.567376
scatter_ratio         168.901775
elongatedness        40.933728
pr.axis_rectangularity 20.582444
max.length_rectangularity 147.998818
scaled_variance      188.631079
scaled_variance.1     439.494076
scaled_radius_of_gyration 174.709716
scaled_radius_of_gyration.1 72.447743
skewness_about        6.364286
skewness_about.1     12.602367
skewness_about.2     188.919527
hollows_ratio        195.632388
class_dup             0.977541
dtype: float64
```

In [10]:

```
veh.drop("class",axis=1,inplace=True)
```

In [11]:

```
meanFiller = lambda x: x.fillna(x.mean())
veh = veh.apply(meanFiller,axis=0)
```

In the above step null values were replaced with means values of there respective columns \ and the class clomn was dropped.


Univariate Analysis

In [12]:

```
sns.distplot(veh["compactness"])
```

Out[12]:

<matplotlib.axes._subplots.AxesSubplot at 0x240748e5588>




In [13]:

```
sns.distplot(veh["circularity"])
```

Out[13]:

<matplotlib.axes._subplots.AxesSubplot at 0x2407500d588>

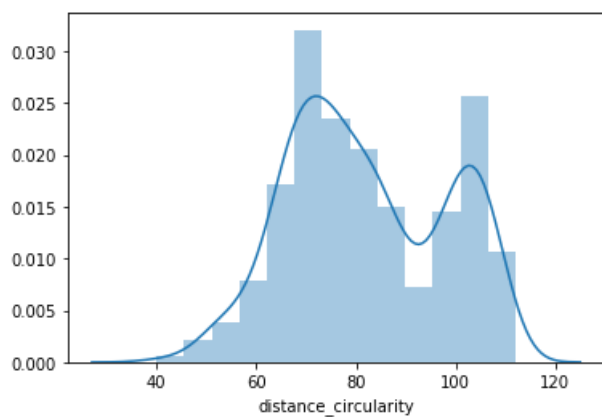


In [14]:

```
sns.distplot(veh["distance_circularity"])
```

Out[14]:

<matplotlib.axes._subplots.AxesSubplot at 0x240750fbf48>




In [15]:

```
sns.distplot(veh["radius_ratio"])
```

Out[15]:

<matplotlib.axes._subplots.AxesSubplot at 0x240751b9c08>




In [16]:

```
sns.distplot(veh["pr.axis_aspect_ratio"])
```

Out[16]:

<matplotlib.axes._subplots.AxesSubplot at 0x24075254048>




In [17]:

```
sns.distplot(veh["max.length_aspect_ratio"])
```

Out[17]:

<matplotlib.axes._subplots.AxesSubplot at 0x240753568c8>





In [18]:

```
sns.distplot(veh["scatter_ratio"])
```

Out[18]:

<matplotlib.axes._subplots.AxesSubplot at 0x2407533d808>



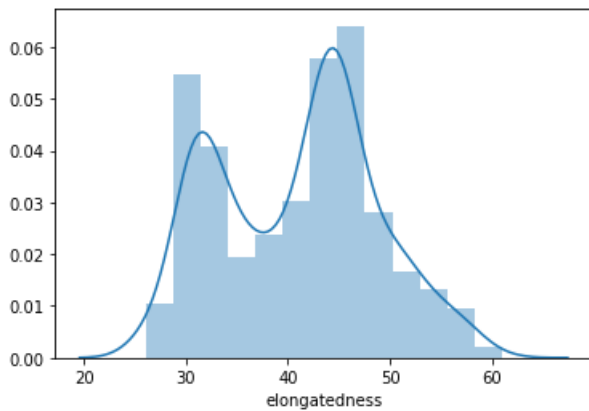


In [19]:

```
sns.distplot(veh["elongatedness"])
```

Out[19]:

<matplotlib.axes._subplots.AxesSubplot at 0x240754d54c8>

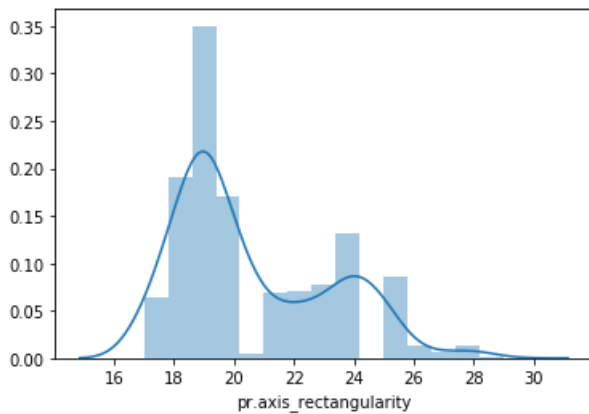


In [20]:

```
sns.distplot(veh["pr.axis_rectangularity"])
```

Out[20]:

<matplotlib.axes._subplots.AxesSubplot at 0x24075564d48>





In [21]:

```
sns.distplot(veh["scaled_variance"])
```

Out[21]:

<matplotlib.axes._subplots.AxesSubplot at 0x240754de708>






In [22]:

```
sns.distplot(veh["scaled_variance.1"])
```

Out[22]:

<matplotlib.axes._subplots.AxesSubplot at 0x240756a1a48>




In [23]:

```
sns.distplot(veh["scaled_radius_of_gyration"])
```

Out[23]:

<matplotlib.axes._subplots.AxesSubplot at 0x2407572b048>





In [24]:

```
sns.distplot(veh["scaled_radius_of_gyration.1"])
```

Out[24]:

<matplotlib.axes._subplots.AxesSubplot at 0x240757c42c8>






In [25]:

```
sns.distplot(veh["skewness_about"])
```

Out[25]:

<matplotlib.axes._subplots.AxesSubplot at 0x2407581a908>




In [26]:

```
sns.distplot(veh["skewness_about.1"])
```

Out[26]:

<matplotlib.axes._subplots.AxesSubplot at 0x2407593d6c8>





In [27]:

```
sns.distplot(veh["skewness_about.2"])
```

Out[27]:

<matplotlib.axes._subplots.AxesSubplot at 0x240769aa088>






In [28]:

```
sns.distplot(veh["hollows_ratio"])
```

Out[28]:

<matplotlib.axes._subplots.AxesSubplot at 0x24076a33688>



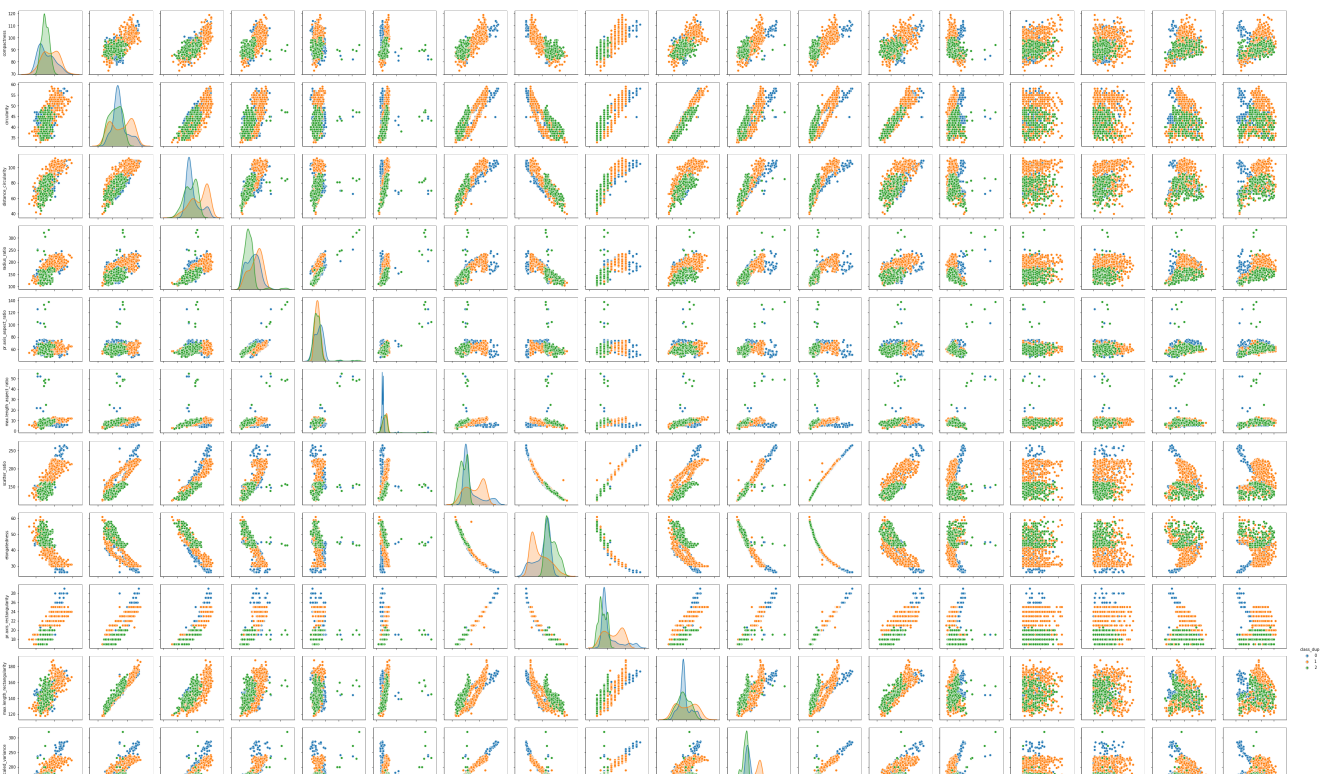
Bivariate Analysis

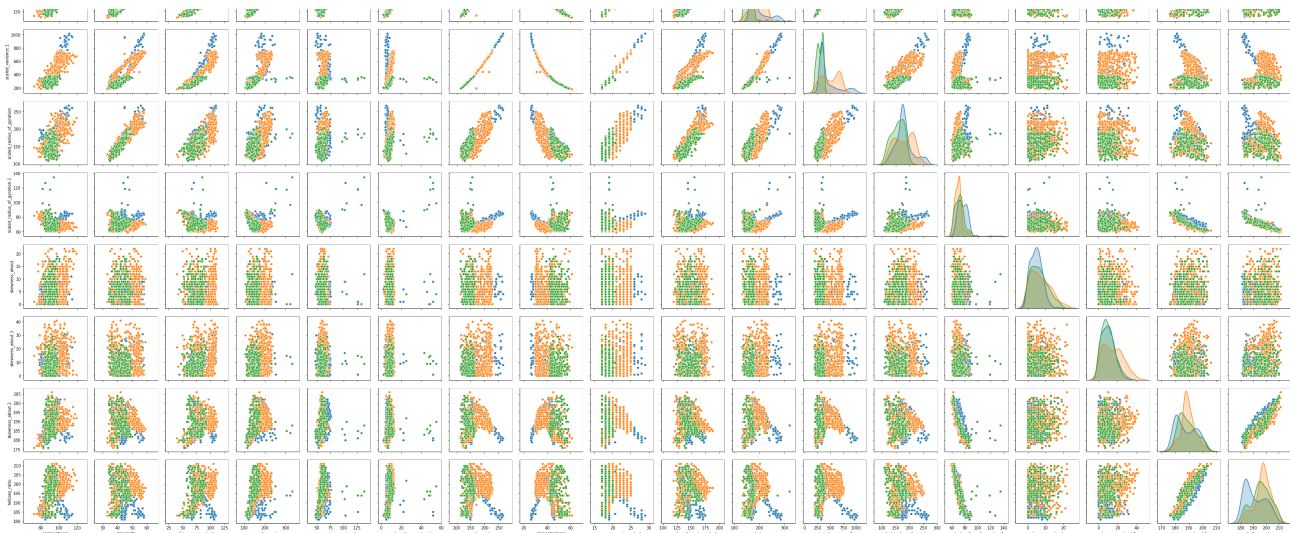
In [29]:

```
sns.pairplot(veh,hue='class_dup',diag_kind='kde')
```

Out[29]:

<seaborn.axisgrid.PairGrid at 0x24076a21888>





Correlation table

In [30]:

```
core=veh.corr()
veh.corr()
```

Out[30]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_aspect_ratio	scatter_ratio
compactness	1.000000	0.685421	0.789909	0.689840	0.091704	0.148249	0.812235
circularity	0.685421	1.000000	0.793016	0.620967	0.153362	0.251208	0.848207
distance_circularity	0.789909	0.793016	1.000000	0.767079	0.158397	0.264550	0.904400
radius_ratio	0.689840	0.620967	0.767079	1.000000	0.663559	0.450036	0.734228
pr.axis_aspect_ratio	0.091704	0.153362	0.158397	0.663559	1.000000	0.648704	0.103715
max.length_aspect_ratio	0.148249	0.251208	0.264550	0.450036	0.648704	1.000000	0.165967
scatter_ratio	0.812235	0.848207	0.904400	0.734228	0.103715	0.165967	1.000000
elongatedness	-0.788643	-0.821901	-0.911435	-0.789795	-0.183264	-0.180041	-0.788643
pr.axis_rectangularity	0.813636	0.844972	0.893128	0.708285	0.079395	0.161592	0.813636
max.length_rectangularity	0.676143	0.961943	0.774669	0.569205	0.127128	0.305943	0.676143
scaled_variance	0.762770	0.796822	0.861980	0.794041	0.273414	0.318955	0.762770
scaled_variance.1	0.815901	0.838525	0.887328	0.720150	0.089620	0.143713	0.815901
scaled_radius_of_gyration	0.585156	0.926888	0.705953	0.536536	0.122111	0.189704	0.585156
scaled_radius_of_gyration.1	-0.250071	0.052642	-0.225852	-0.180819	0.152776	0.295574	-0.250071
skewness_about	0.235687	0.144394	0.113813	0.048720	-0.058481	0.015439	0.235687
skewness_about.1	0.157387	-0.011851	0.265553	0.173832	-0.032134	0.043489	0.157387
skewness_about.2	0.298526	-0.105645	0.145563	0.382129	0.239849	-0.026180	0.298526
hollows_ratio	0.365552	0.045318	0.332095	0.471262	0.267724	0.143919	0.365552
class_dup	-0.033796	-0.159804	-0.064902	-0.182270	-0.098318	0.207619	-0.033796

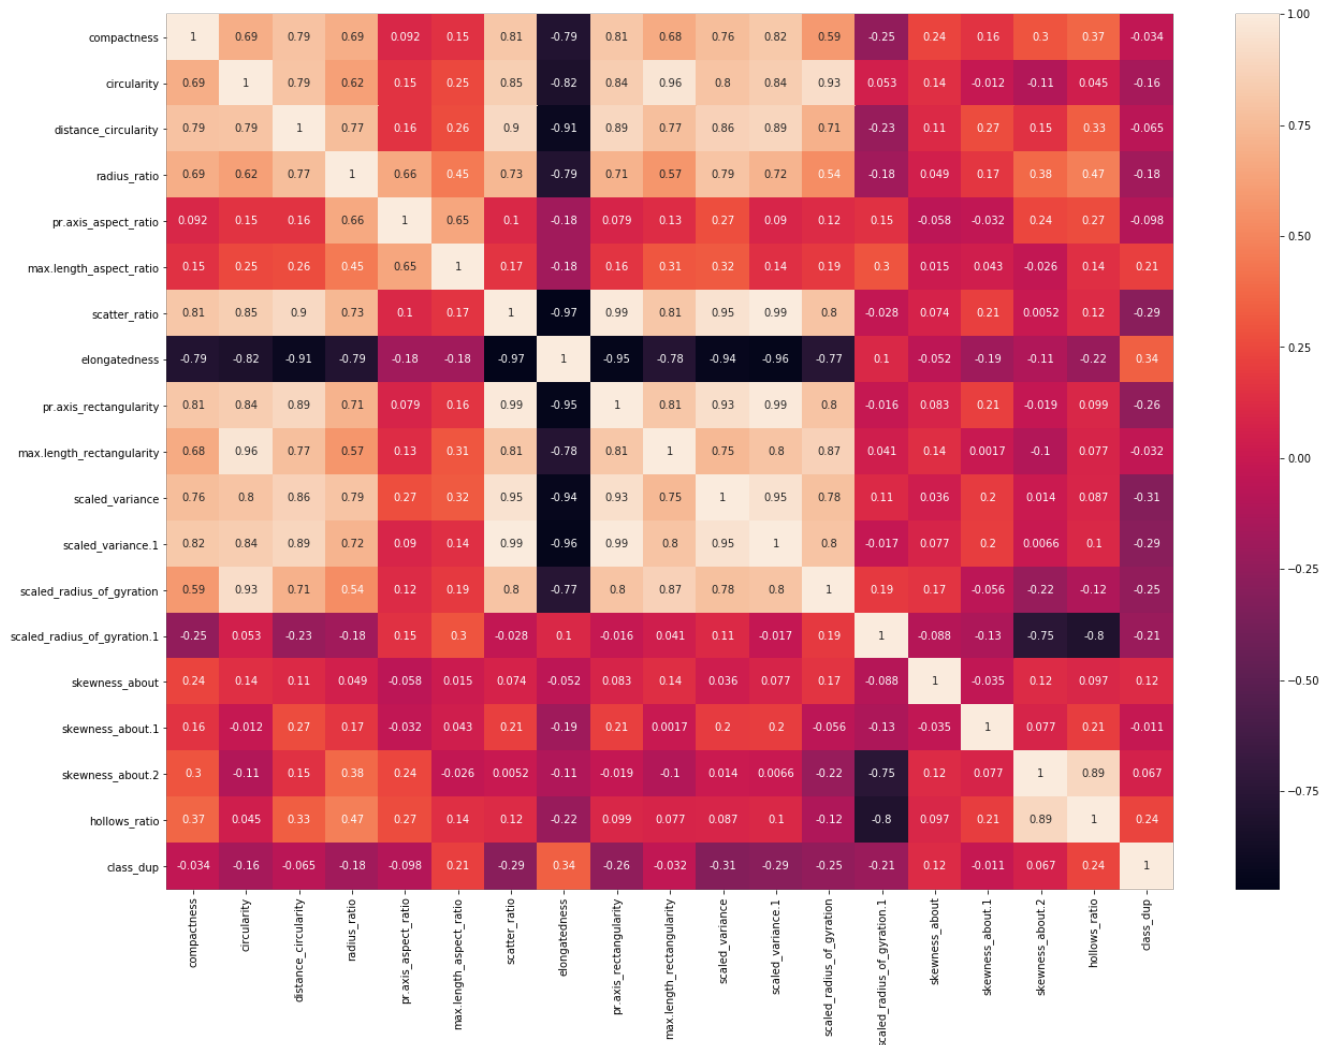
Correlation Heatmap

In [31]:

```
plt.figure(figsize = (21,15))
sns.heatmap(core,annot=True)
```

Out[31]:

<matplotlib.axes._subplots.AxesSubplot at 0x24004a91288>



From the bivariate analysis, correlation table and correlation heatmap.\ The columns the is of least significance to target variable was found.

Preprocessing Data for analysis

In [32]:

```
veh.drop("compactness",axis=1,inplace=True)
veh.drop("distance_circularity",axis=1,inplace=True)
veh.drop("pr.axis_aspect_ratio",axis=1,inplace=True)
veh.drop("max.length_rectangularity",axis=1,inplace=True)
veh.drop("scatter_ratio",axis=1,inplace=True)
veh.drop("pr.axis_rectangularity",axis=1,inplace=True)
veh.drop("skewness_about.2",axis=1,inplace=True)
```

In [33]:

```
X= veh.drop(["class_dup"],axis=1,)
Y= veh[["class_dup"]]
```

In [34]:

```
XScaled = X.apply(zscore)
```

Splitting the data into test,train

In [35]:

```
x_train, x_test, y_train, y_test = train_test_split(XScaled, Y, test_size=0.3, random_state=7)
```

SVM Analysis

In [36]:

```
clf = SVC(C=.1, kernel='linear', gamma= 1)
```

In [37]:

```
clf.fit(x_train , y_train.values.ravel())
```

Out[37]:

```
SVC(C=0.1, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=1, kernel='linear',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

In [38]:

```
y_pred = clf.predict(x_test)
```

In [39]:

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.8582677165354331

In [40]:

```
clf = SVC(C=.1, kernel='rbf', gamma= 1)
clf.fit(x_train , y_train.values.ravel())
```

Out[40]:

```
SVC(C=0.1, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=1, kernel='rbf', max_iter=-1,
    probability=False, random_state=None, shrinking=True, tol=0.001,
    verbose=False)
```

In [41]:

```
y_pred = clf.predict(x_test)
```

In [42]:

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.531496062992126

In [43]:

```
clf = SVC(C=.1, kernel='poly', gamma= 1)
clf.fit(x_train , y_train.values.ravel())
```

Out[43]:

```
SVC(C=0.1, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=1, kernel='poly',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

In [44]:

```
y_pred = clf.predict(x_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.9133858267716536

In [45]:

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Confusion Matrix:


```
[[ 55   7   4]
 [  3 120   4]
 [  1   3  57]]
```

In [46]:

```
knnm=metrics.confusion_matrix(y_test,y_pred)
knn_m = pd.DataFrame(knnm, index = [i for i in ["2","1","0"]], columns = [i for i in ["Predict 2","Predict 1","Predict 0"]])
plt.figure(figsize = (7,5))
sns.heatmap(knn_m, annot=True)
```

Out[46]:

<matplotlib.axes._subplots.AxesSubplot at 0x2400405fa88>



In [47]:

```
print("Classification Report")
print(metrics.classification_report(y_test,y_pred, labels=[2,1, 0]))
```

Classification Report				
	precision	recall	f1-score	support
2	0.88	0.93	0.90	61
1	0.92	0.94	0.93	127
0	0.93	0.83	0.88	66
accuracy			0.91	254
macro avg	0.91	0.90	0.91	254
weighted avg	0.91	0.91	0.91	254

In [48]:

```
clf = SVC(C= .1, kernel='sigmoid', gamma= 1)
clf.fit(x_train , y_train.values.ravel())
```

Out[48]:

```
SVC(C=0.1, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=1, kernel='sigmoid',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

In [49]:

```
y_pred = clf.predict(x_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.5236220472440944

After trying the SVM algorithm in the different kernels. Following inference was made:-

- The highest model accuracy is when the algorithm is running in a polynomial kernel.
- The highest accuracy score was 91.
- Confusion matrix and classification reports were found for the same.

K-fold and cross-validation

In [50]:

```
kfold = KFold(n_splits=10, random_state=7, shuffle=True)
clf= SVC(C= .1, kernel='poly', gamma= 1)
model=clf.fit(x_train , y_train.values.ravel())
results = cross_val_score(model, XScaled, Y.values.ravel() , cv=kfold)
print(results)
print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

```
[0.94117647 0.91764706 0.90588235 0.95294118 0.96470588 0.83529412
 0.92857143 0.85714286 0.86904762 0.91666667]
Accuracy: 90.891% (4.039%)
```

After using K-fold validation.

- The range of accuracy of the model was from 86% to 94% at a 95% confidence level.

Identifying Eigen vectors using PCA

In [51]:

```
covMatrix = np.cov(XScaled, rowvar=False)
print(covMatrix)
```

```
[ [ 1.00118343  0.62170187  0.25150523 -0.82287347  0.7977645  0.83951746
   0.92798524  0.05270433  0.14456452 -0.01186527  0.04537164]
 [ 0.62170187  1.00118343  0.45056889 -0.79072934  0.79498064  0.72100219
   0.5371705  -0.1810333  0.04877731  0.17403799  0.47181974]
 [ 0.25150523  0.45056889  1.00118343 -0.18025396  0.31933203  0.14388275
   0.18992805  0.29592367  0.01545721  0.04354026  0.14408905]
 [-0.82287347 -0.79072934 -0.18025396  1.00118343 -0.93782312 -0.95620413
  -0.76693543  0.10360323 -0.05205875 -0.18591103 -0.21697531]
 [ 0.7977645  0.79498064  0.31933203 -0.93782312  1.00118343  0.94814159
   0.77989661  0.11243163  0.03604752  0.19549063  0.08669654]
 [ 0.83951746  0.72100219  0.14388275 -0.95620413  0.94814159  1.00118343
   0.79701241 -0.0166278  0.07706469  0.20181158  0.10388468]
 [ 0.92798524  0.5371705  0.18992805 -0.76693543  0.77989661  0.79701241
   1.00118343  0.19166642  0.16656805 -0.05603902 -0.1182971 ]
 [ 0.05270433 -0.1810333  0.29592367  0.10360323  0.11243163 -0.0166278
```

```

0.19166642 1.00118343 -0.08840848 -0.12656621 -0.8035581 ]
[ 0.14456452 0.04877731 0.01545721 -0.05205875 0.03604752 0.07706469
 0.16656805 -0.08840848 1.00118343 -0.03506456 0.09698477]
[-0.01186527 0.17403799 0.04354026 -0.18591103 0.19549063 0.20181158
-0.05603902 -0.12656621 -0.03506456 1.00118343 0.20533271]
[ 0.04537164 0.47181974 0.14408905 -0.21697531 0.08669654 0.10388468
-0.1182971 -0.8035581 0.09698477 0.20533271 1.00118343]]

```

In [52]:

```

pca = PCA(n_components=6)
pca.fit(XScaled)

```

Out[52]:

```

PCA(copy=True, iterated_power='auto', n_components=6, random_state=None,
     svd_solver='auto', tol=0.0, whiten=False)

```

In [53]:

```

print(pca.explained_variance_)

```

```

[5.18662721 2.05550394 1.17384819 1.05528084 0.88494564 0.32133759]


```

In [54]:

```

plt.bar(list(range(1,7)),pca.explained_variance_ratio_,alpha=0.5, align='center')
plt.ylabel('Variation explained')
plt.xlabel('eigen Value')
plt.show()

```



In [55]:

```

pca3 = PCA(n_components=6)
pca3.fit(XScaled)
print(pca3.components_)
print(pca3.explained_variance_ratio_)
Xpca3 = pca3.transform(XScaled)

```

```

[[ 0.39611623 0.36809097 0.14874228 -0.42089468 0.42119261 0.41767613
  0.37539781 -0.00394397 0.05149428 0.07315524 0.086749 ]
 [ 0.1160752 -0.20374372 0.07430875 0.05260597 0.06271017 0.01894014
  0.22546146 0.6382345 -0.06066981 -0.23779087 -0.64345404]
 [-0.17512039 0.22052399 0.65318935 0.06155997 0.07617988 -0.09618641
 -0.22467436 0.27090577 -0.45675851 0.36913349 0.08027758]
 [ 0.04984377 0.14037895 0.56535966 0.14808997 -0.09600545 -0.18739837
  0.02838387 0.05231413 0.59254708 -0.44522476 0.19756629]
 [-0.07623596 -0.1169546 -0.01159183 0.03607123 0.02585673 0.0410813
 -0.02828491 0.16883607 0.63259572 0.72603927 -0.14178125]
 [-0.52481012 0.39545311 -0.22829585 -0.19894854 0.31027591 0.17083502
 -0.45312434 0.2015486 0.17776456 -0.24707486 -0.11156243]]
[0.47095422 0.18664312 0.10658733 0.09582122 0.08035451 0.02917798]

```

In [56]:

```
Xpca3
```

Out[56]:

```
array([[ 0.12968519, -0.34629014,  0.22002525,  0.17720568,  0.04526196,
        -0.76880923],
       [-1.38405839, -0.43238023, -0.01926176,  0.47611585,  0.51626564,
        -0.19116416],
       [ 2.80683385,  0.23658889, -0.92230398,  1.06053268,  0.48296298,
        0.47390012],
       ...,
       [ 3.3245831 , -0.28744424,  0.18466963,  0.26263644, -1.57503553,
        0.23190269],
       [-2.49331101, -1.01044899,  1.01014057, -1.37892084,  0.2449125 ,
        -0.53232961],
       [-3.77918478,  0.18938257,  0.55137633, -1.1286787 ,  0.16053356,
        -0.33219449]])
```

In [57]:

```
Xpca3.shape
```

Out[57]:


```
(846, 6)
```


In [58]:

```
sns.pairplot(pd.DataFrame(Xpca3))
```

Out[58]:

```
<seaborn.axisgrid.PairGrid at 0x240050e1c48>
```





In [59]:

```
x_train, x_test, y_train, y_test = train_test_split(Xpca3, Y, test_size=0.3, random_state=7)
```

SVM Analysis using the eigen vectors

In [60]:

```
clf = SVC(C=.1, kernel='linear', gamma=1)
```

In [61]:

```
clf.fit(x_train, y_train.values.ravel())
```

Out[61]:

```
SVC(C=0.1, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=1, kernel='linear',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

In [62]:

```
y_pred = clf.predict(x_test)
```

In [63]:

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.8110236220472441

In [64]:

```
clf = SVC(C=.1, kernel='rbf', gamma=1)
clf.fit(x_train, y_train.values.ravel())
y_pred = clf.predict(x_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.5511811023622047

In [65]:

```
clf = SVC(C=.1, kernel='poly', gamma=1)
clf.fit(x_train, y_train.values.ravel())
y_pred = clf.predict(x_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.8346456692913385

In [66]:

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Confusion Matrix:

```
[[ 42  20   4]
 [   6 116   5]
 [   1   6  54]]
```

In [67]:

```
knnm=metrics.confusion_matrix(y_test,y_pred)
knn_m = pd.DataFrame(knnm, index = [i for i in ["2","1","0"]],columns = [i for i in ["Predict2","Predict 1","Predict 0"]])
plt.figure(figsize = (7,5))
sns.heatmap(knn_m, annot=True)
```

Out[67]:

<matplotlib.axes._subplots.AxesSubplot at 0x2400c2d8848>



In [68]:

```
print("Classification Report")
print(metrics.classification_report(y_test,y_pred, labels=[2,1, 0]))
```

```
Classification Report
              precision    recall  f1-score   support

     2       0.86      0.89      0.87         61
     1       0.82      0.91      0.86        127
     0       0.86      0.64      0.73         66

 accuracy      0.83         254
 macro avg     0.84         254
weighted avg     0.84         254
```

In [69]:

```
clf = SVC(C=.1, kernel='sigmoid', gamma= 1)
clf.fit(x_train , y_train.values.ravel())
y_pred = clf.predict(x_test)
print("Accuracy:",accuracy_score(y_test, y_pred))
```

Accuracy: 0.5118110236220472

After running the SVM model with eigenvectors in different kernels. The following was inferred.

- The maximum accuracy was again found in the polynomial kernel.
- The accuracy of this model was 83.
- A drop of around 6% was found but this expected as the columns are reduced

K-fold validation with PCA value

In [70]:

```
kfold = KFold(n_splits=10, random_state=7, shuffle=True)
clf= SVC(C= .1, kernel='poly', gamma= 1)
model=clf.fit(x_train , y_train.values.ravel())
results = cross_val_score(model, Xpca3, Y.values.ravel() , cv=kfold)
print(results)
print("Accuracy: %.3f%% (%.3f%%)" % (results.mean()*100.0, results.std()*100.0))
```

```
[0.85882353 0.87058824 0.81176471 0.84705882 0.84705882 0.75294118
 0.80952381 0.82142857 0.82142857 0.89285714]
Accuracy: 83.335% (3.706%)
```

After using K-fold validation with eigen vectors.

- The range of accuracy of the model was from 80% to 87% at 95% confidence level.

Even after only selecting six column that covered 95% of data using PCA.

- The accuracy dropped from 91 to 83
- This is result was within the expected region.